

Prometheus *a practical workshop* v1.1

Tomer Gabel
@ JCON Slovenia
May 2025



Java Edition

The Pillars of Observability

1. Logs

2. Metrics

3. Traces

The Pillars of Observability

~~1. Logs~~

2. Metrics

~~3. Traces~~

Our focus today.

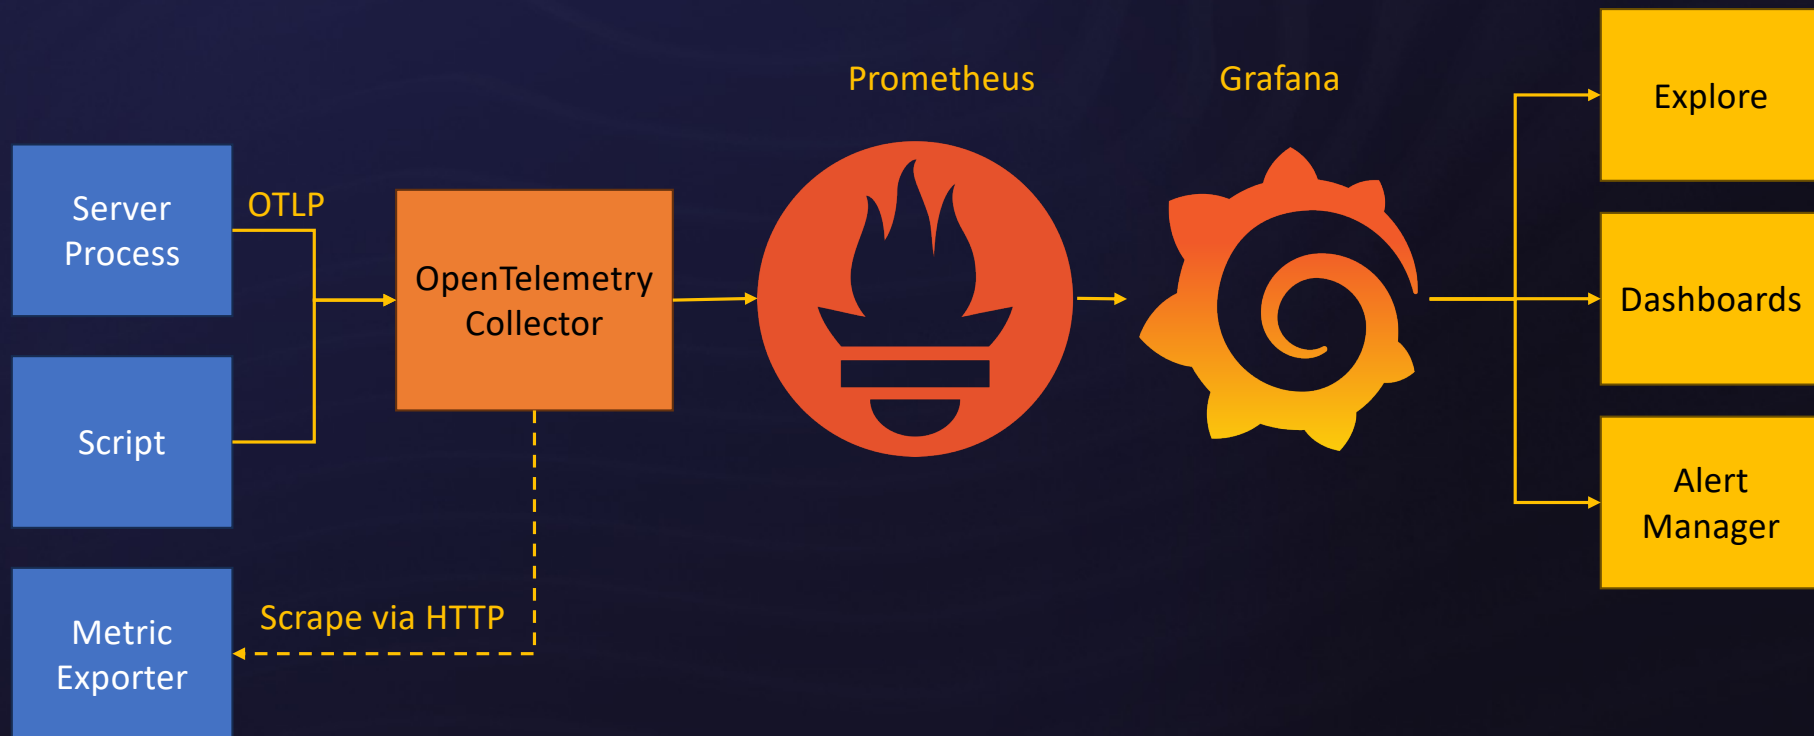


Metrics:

Answer *quantitative* questions

1. *How many* requests did I get?
2. *How long* did they take to process?
3. *How much* memory am I using?
4. *How much* free space is left on the disk?

Enter: Prometheus



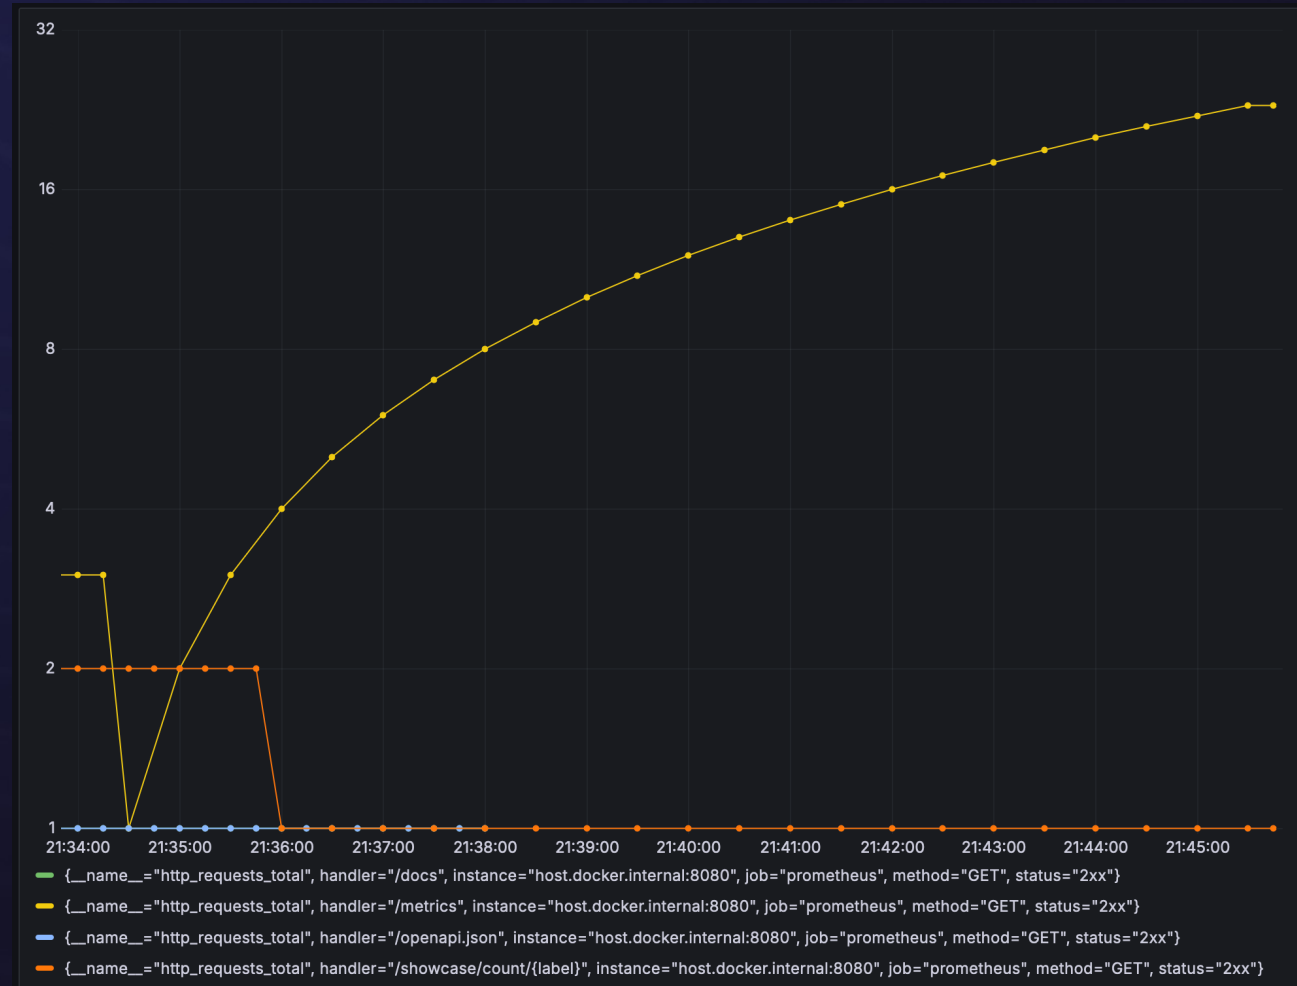
Metrics 101

1. Name

2. Labels

3. Time

4. Value



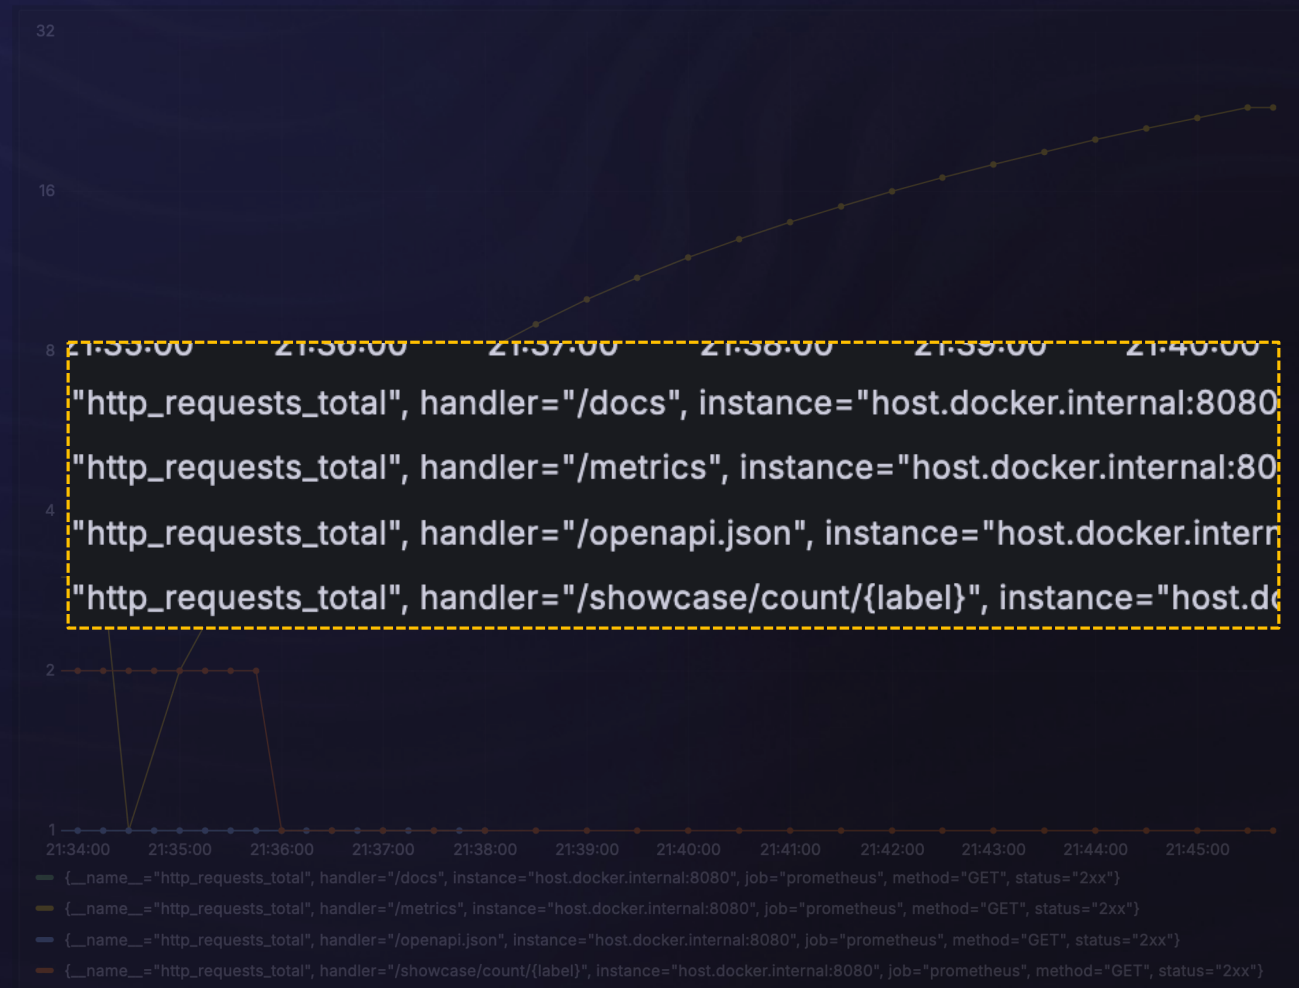
Metrics 101

1. Name

2. Labels

3. Time

4. Value



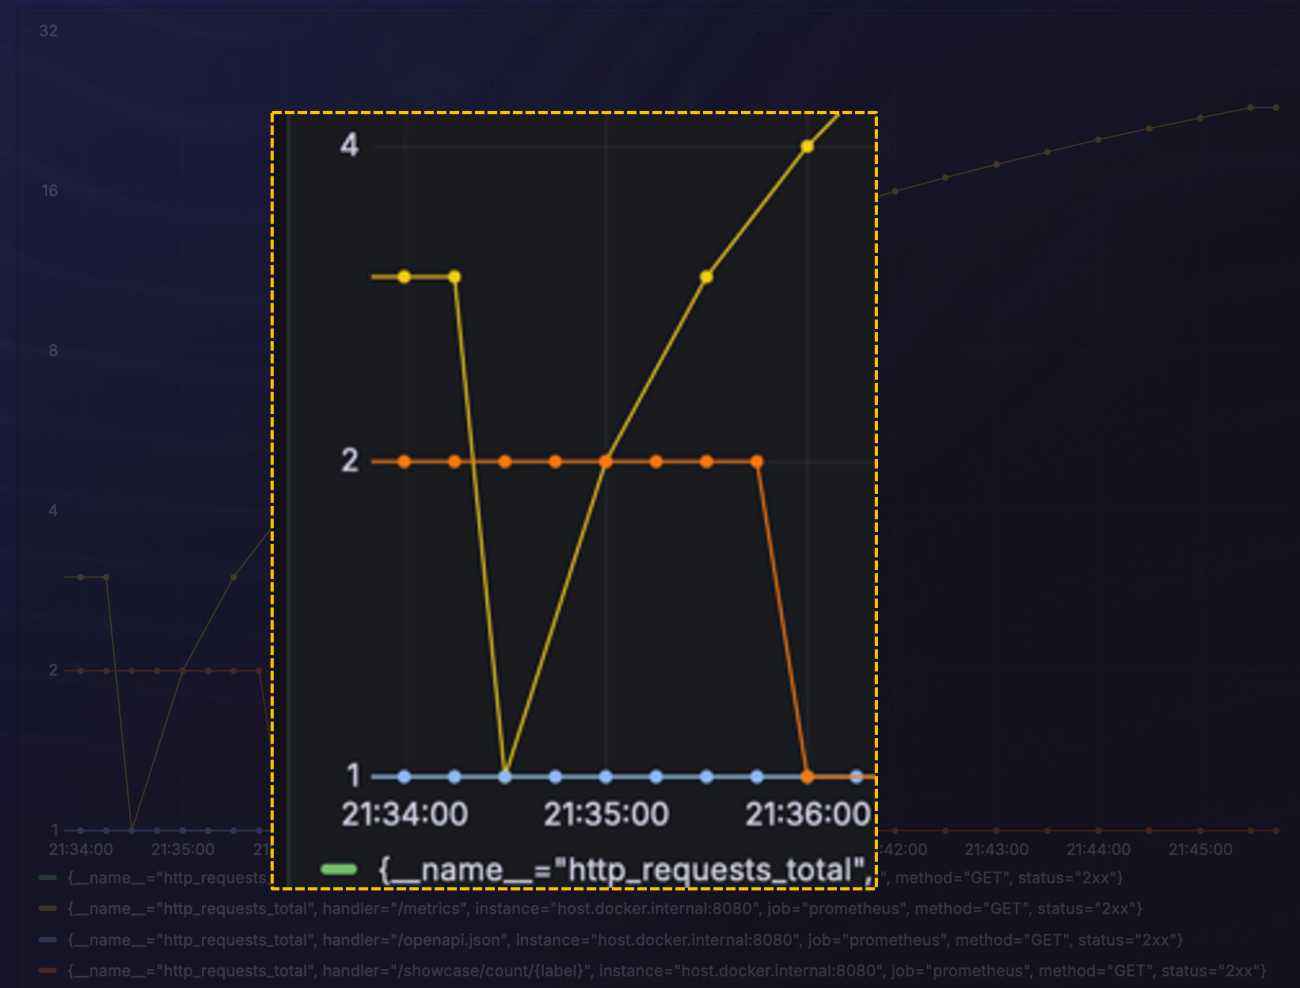
Metrics 101

1. Name

2. Labels

3. Time

4. Value



A word on cardinality

1. Prometheus is analogous to a *flat database*
 - Partitioned by *time*
 - Labels are like *indexed columns*
2. The DB size is therefore:
 - $\text{Timeframe} * \# \text{ metrics} * \# \text{ label values}$

Prometheus metric types: Counter


1. As the name implies, a metric that *counts*
2. Can go up or down by any amount
3. *Stateful* – value managed by *Prometheus*
4. Examples:
 - # of server requests (by method, path, ...)
 - # of events ("times a user logged in")

Counting the Java way

```
OpenTelemetry otel = ...;
```

```
var meter = otel.getMeter("controller");  
var counter = meter  
    .counterBuilder("my_count")  
    .setDescription("Event count")  
    .build();
```

```
counter.add(  
    1,  
    Attributes.of(stringKey("my_label"), "value")  
);
```



Prometheus metric types: Gauge

1. A metric that is sampled *on demand*
2. Each sample is *independent*
3. Common examples:
 - CPU core temperature in °C
 - Free space on /dev/sda2


Gauging the Java way

```
OpenTelemetry otel = ...;
```

```
var meter = otel.getMeter("controller");
```

```
var gauge = meter  
    .gaugeBuilder("my_value")  
    .setDescription("Some stateful value")  
    .ofLongs()  
    .build();
```

```
gauge.set(15);
```



Prometheus metric types: Histograms


1. Some metrics *cannot be represented* with one value
 - “*How long* do requests to /login take?”
 - Do you mean average? Mean? 90th percentile?
2. These values produce a *distribution*
3. Managed with buckets. Latency, for example:
 - 0-10ms, 10-100ms, 100-1000ms, ...

Histograms the Java way

```
OpenTelemetry otel = ...;
```

```
var meter = otel.getMeter("controller");  
var histogram = meter  
    .histogramBuilder("my_duration")  
    .setDescription("Event duration")  
    .setUnit("ms")  
    .ofLongs()  
    .build();
```

```
histogram.record(3400);
```



Not all histograms are made equal

Latencies vary wildly (ms vs second), as do *sizes* (payload in KB, disk in GB).

Views to the rescue!

```
var buckets = List.of(0.0, 500.0, 1000.0, 10000.0);
var meterProvider = SdkMeterProvider.builder()
    .registerView(
        InstrumentSelector.builder().setName("my_duration").build(),
        View.builder()
            .setAggregation(Aggregation.explicitBucketHistogram(buckets))
            .build()
    ).build();
```


The background is a dark navy blue. It features a series of concentric circles on the right side, which transition into a spiral pattern on the left side. The spiral is composed of many thin, overlapping lines that create a sense of depth and movement. The text 'Lab Time' is centered in the middle of the image, written in a bold, yellow, sans-serif font.

Lab Time

Showcase

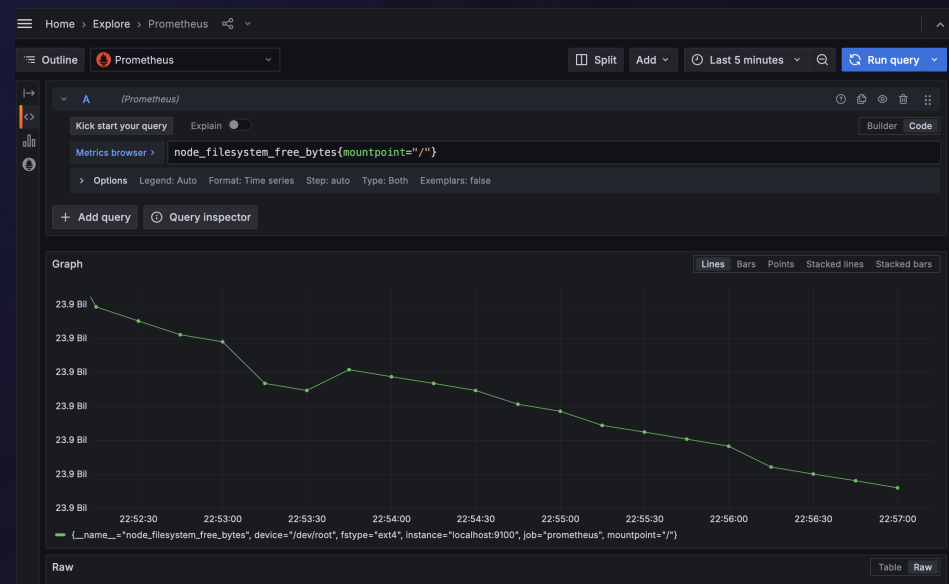
1. *Familiarize yourself* with the lab setup
2. *Get handsy* with Grafana
3. Open lab-showcase from the *class materials*
4. You have *30 minutes* to explore!

The background is a dark navy blue. It features a large, faint, light-blue graphic on the left side. This graphic consists of a series of concentric circles that form a spiral, resembling a nautilus shell or a stylized eye. The spiral starts from the center and moves outwards towards the top-left corner of the frame.

Part 2: PromQL

PromQL: Fear, Terror and Ruthless Efficiency

1. Extensive, powerful
2. ... not entirely trivial



PromQL: Fear, Terror and Ruthless Efficiency

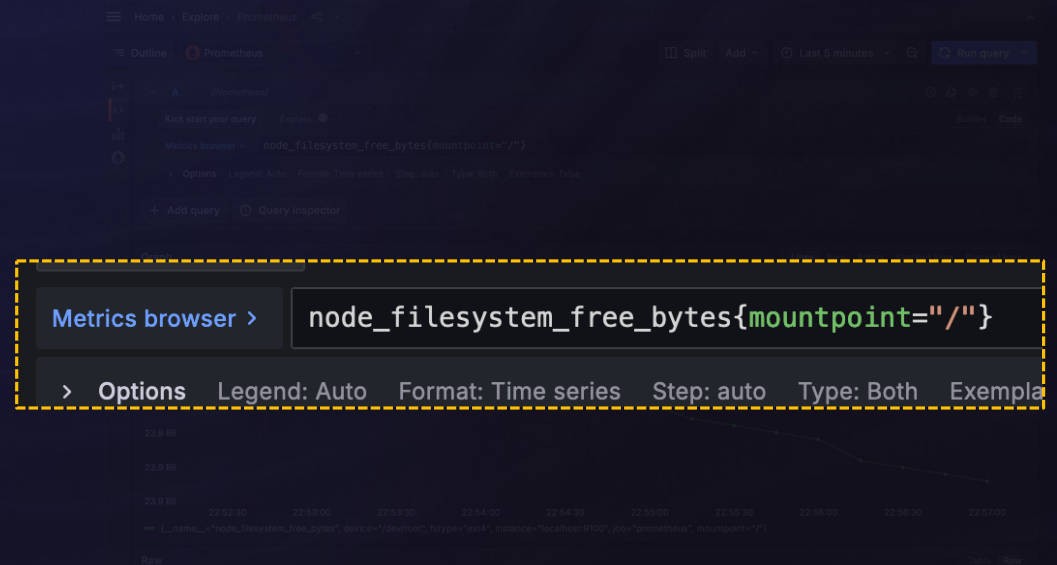
1. Extensive, powerful

2. ... not entirely trivial

3. Basically:

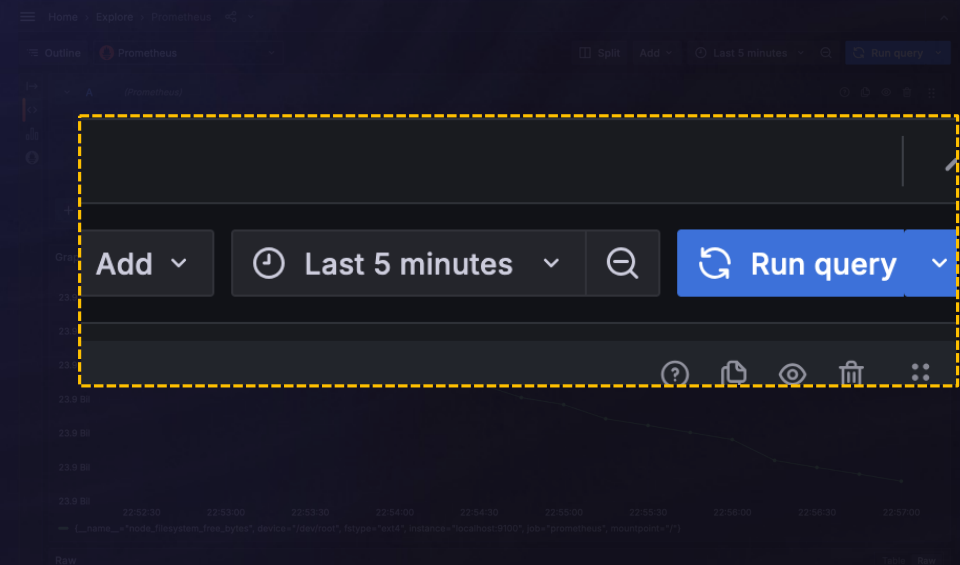
- A *query*

- A *timeframe*



PromQL: Fear, Terror and Ruthless Efficiency


1. Extensive, powerful
2. ... not entirely trivial
3. Basically:
 - A *query*
 - A *timeframe*



PromQL: Maths!

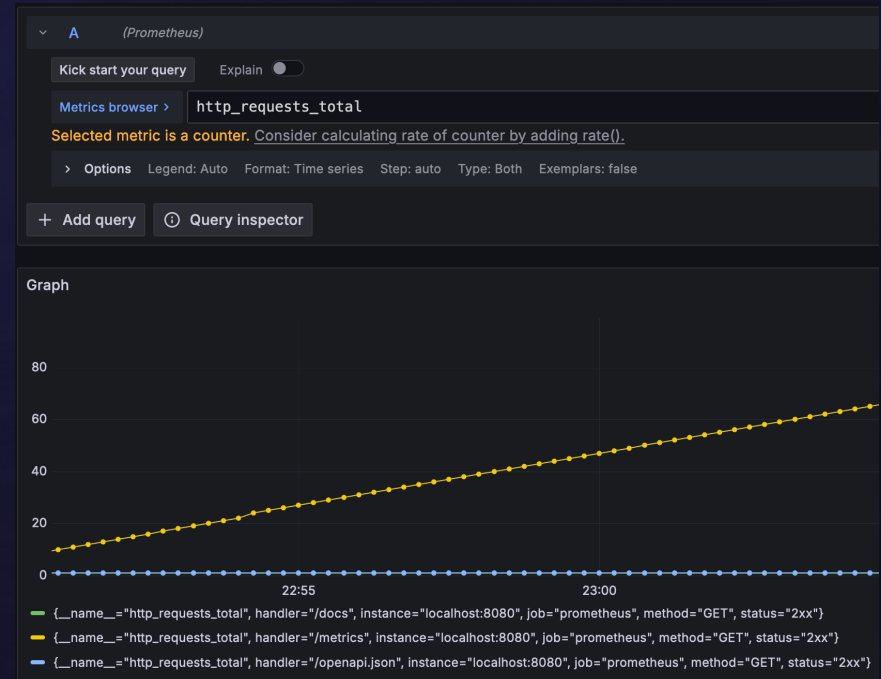
1. [Metric, labels, time] yields **one value**
2. Some observations require **multiple** values
3. Expressed with basic arithmetic:
 - % free disk space = **free** disk bytes / **total** disk bytes
 - % CPU = **CPU used seconds** / (# seconds * **# cores**)

PromQL: Maths



PromQL: Rates

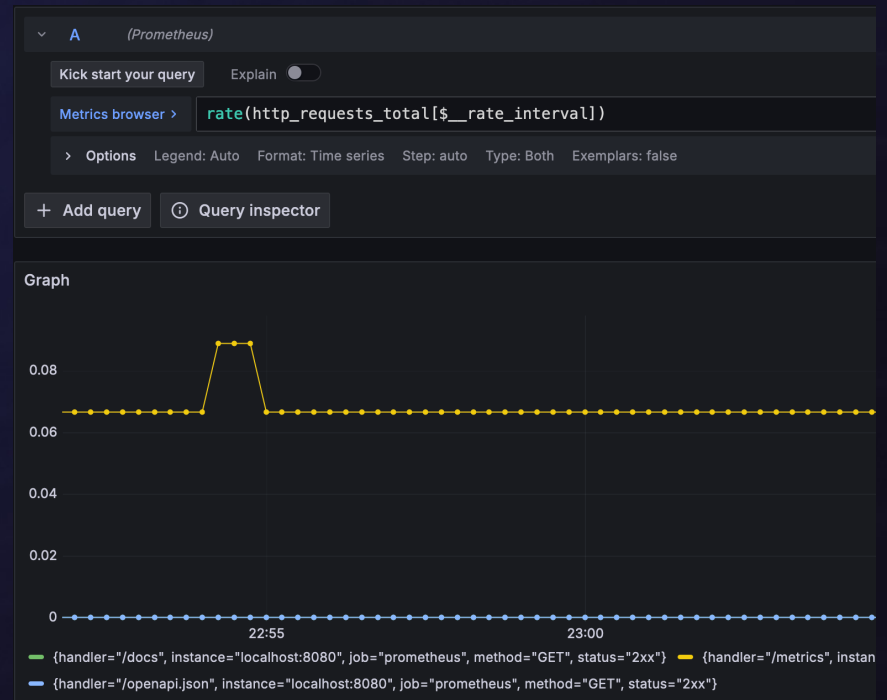
1. Counters are very useful
2. Can't use them *naïvely*



PromQL: Rates

1. Counters are very useful
2. Can't use them *naïvely*
3. We need a *rate derivation*
4. Magic incantation for the win:

`rate(metric{...}[$__rate_interval])`



PromQL: Aggregations

1. Consider *http_server_duration_milliseconds_count*
 - A simple counter
 - Potentially many services, paths
2. Suppose you want the grand total?
`sum(http_server_duration_milliseconds_count)`

PromQL: Aggregations

1. Consider *http_server_duration_milliseconds_count*
 - A simple counter
 - Potentially many services, paths
2. How about the 5 most used paths?
`topk(5, http_server_duration_milliseconds_count)`

PromQL: Aggregations

1. What are the top 3 most time-consuming endpoints?

```
topk(3, http_server_duration_milliseconds_sum)
```

2. How many processes run which Python version?

```
count by (major) (python_info)
```

3. What is the highest rate of pagefaults recorded?

```
max(rate(node_vmstat_pgfault[$__rate_interval]))
```


A dark blue background with a white line that starts at the top left, slopes down to a minimum point, and then slopes back up towards the right edge, creating a valley-like shape.

And now,
what you've all
been waiting for

Quantiles (a.k.a. percentiles)


1. A histogram is a set of *counters*
 - 0-100ms, 100-200ms, 200-300ms...
2. What is the 95th *percentile*?
 - Linear regression FTW... but maths!
3. PromQL to the rescue:

`histogram_quantile(0.99, http_server_duration_milliseconds_bucket)`




Graph

Lines Bars Points **Stacked lines** Stacked bars



Graph

Lines Bars Points **Stacked lines** Stacked bars



Key Takeaways

1. Instrument early, instrument often

There's no such thing as "too much data"

2. Beware high *cardinality*

Aggregate, don't transact (that's what logs/traces are for)

3. Learn ye *PromQL*

That's where the real leverage is

The background is a dark navy blue. It features a series of concentric circles that are slightly offset, creating a sense of depth. Overlaid on these circles is a spiral pattern that starts from the center and moves outwards, also appearing to have a 3D effect with varying shades of blue. The text 'Lab time, again!' is centered in a bold, yellow, sans-serif font.

Lab time, again!



tomer@substrate.co.il



@tomerg



<https://github.com/holograph/prometheus-workshop-service-python>

Thank you for your attention

Questions?

